

Guidelight MVP Sprint Plan

A detailed playbook for building the first working version of Guidelight in one focused day — or across several smaller focused sessions

- > **Purpose of this document**
- > This is a long-form, detailed plan for how you (Justin) can spend a focused day building the **Guidelight** MVP using **Cursor**, **GPT-5.1** , **Claude Sonnet 4.5** , and **Codex IDE**.
- > It's meant to be something you can **read tonight** to get mentally ready, and **glance at tomorrow** as a guide while you work.

How to Use This Plan

This playbook was originally written as a one-day sprint, but you can also use it as a sequence of small work sessions.

- When it says “morning”, “today”, or “end of day”, read that as “the next time I sit down to work on Guidelight” or “the end of this work session”.
- It’s okay to pause between phases and pick up on another day; you still follow the same overall order.
- The self-care notes (water, food, breaks, wrap-up) are meant for any session, not just one big marathon.

Use whatever chunks of it fit the time and energy you have in real life.

0. Big Picture

0.1 What “MVP” means for Guidelight

By the end of the sprint, the **Guidelight MVP** should:

- **Boot and load in a browser** on your dev machine (Vite + React + TS).
- Connect to a **dev Supabase/Postgres** instance (read and write).
- Have a functional **Staff View** where a budtender (or you) can:
 - See the list of budtender picks.
 - Create a new pick.
 - Edit an existing pick.
 - Delete a pick or archive it (even if this is basic / rough).
- Have a simple, clean **Customer View** where:
 - Picks are displayed grouped or tagged by budtender.
 - Info is clear and legible, with only the most important details (name, strain/product, 1–2 key attributes).
- Have **docs that match reality** at a high level:
 - `README.md` roughly describes what’s actually there.
 - `GUIDELIGHT_SPEC.md` is mostly correct (minor future work is ok).
 - `ARCHITECTURE_OVERVIEW.md` reflects the structure you actually implemented.

Not perfect. Not pretty. But honest, working, and understandable.

1. Tooling & Model Strategy (Reference)

This is the **fixed strategy** we are committing to for this sprint.

1.1 Environment & tools

- **IDE:** Cursor on Windows (your main home).
- **Backend:** Supabase/Postgres (dev project).
- **Frontend stack:** Vite + React + TypeScript.
- **Data fetching:** Initially, simple Supabase client functions and/or a thin API layer in `src/lib/api`. React Query can be introduced later if needed, but it's not required for day one.

1.2 Models and their roles

You will **manually choose the model in Cursor**. The agent cannot do that itself.

- **Primary model – GPT-5.1 Thinking (in Cursor)**
 - Use for:
 - Most feature implementation.
 - Data layer wiring.
 - UI components & layout.
 - Small and medium refactors.
 - Think of it as your **default co-dev**.
- **Specialist model – Claude Sonnet 4.5 (in Cursor)**
 - Use **only when you choose to switch the model** for:
 - Large, cross-cutting refactors (e.g. cleaning up multiple views + data layer).
 - Deep, multi-file debugging when things feel tangled.
 - Architecture review / “make this section clearer and simpler” tasks.
 - Think of it as your **surgical refactor/debug expert**.
- **Composer model (Cursor's own)**
 - Use only when you explicitly want **fast scaffolding or experiments**:
 - Rough component layouts you intend to clean up.
 - Quick “what if” variants.
 - Not for final code that you intend to keep unchanged.
- **Cursor Auto mode**
 - **Do not** use Auto for this repo during the sprint.
 - Always explicitly set a model for Guidelight work.
- **Codex IDE**
 - Use near the **end of the day** as a **reviewer**, not as your main implementer.
 - It should:

- Review diffs / PR.
- Point out type issues, runtime risks, and complexity.
- Suggest specific fixes that you then implement with GPT-5.1 in Cursor.

1.3 The dev agent file

- `GUIDELIGHT_DEV_AGENT.md` is your **AI configuration / system prompt** for Guidelight inside Cursor.
- You will:
 - Keep it in the repo, and
 - Paste its contents into a **custom Cursor agent** for this project.
- The file now instructs the model to:
 - Respect project docs.
 - Use MCP tools for schema and API, not guess.
 - Suggest which model (GPT-5.1 vs Claude) would be best for a task, while still answering as the currently selected model.

You can think of `GUIDELIGHT_DEV_AGENT.md` as the **job description and rulebook** for any AI that works on this repo.

2. Overall Day Structure

Here's the **high-level flow** of the day:

1. **Morning Setup & Intent** – Get machine/tools ready & set mindset.
2. **AI Reads Docs & Proposes a Plan** – GPT-5.1 internalizes the repo and outlines steps.
3. **App Shell & Routing** – Basic Vite app + Customer/Staff routes.
4. **Data Layer & Supabase Wiring** – Typed API functions, real data.
5. **Staff View MVP** – CRUD or near-CRUD for picks.
6. **Customer View MVP** – Display picks with clean layout.
7. **Refactor Pass with Claude** – One controlled refactor sweep.
8. **Codex Review & Docs Sync** – External review, fix issues, update docs.
9. **Wrap-Up & Notes** – Capture what's done and what's next.

We'll expand each phase below with **recommended prompts, mental notes, and breaks**.

3. Phase-by-Phase Plan

Phase 0 – Morning Reset & Priming (15–30 minutes)

Goal: Start clear and grounded, not frantic.

Steps:

1. **Hydrate & food**

- Drink water. Have coffee/tea if you'd like.
 - Eat *something* (even if small) to avoid brain fog mid-morning.
2. **Light review away from the computer (optional but helpful)**
- On your phone or tablet, skim:
 - `GUIDELIGHT_SPEC.md`
 - `ARCHITECTURE_OVERVIEW.md`
 - Just to recall:
 - Who uses Guidelight.
 - What the two main views are supposed to do.
 - The data model at a high level.

3. **Set a simple intention**
- A sentence in your head, on a sticky note, or in a small text file:
 - “Today I’m building a functional, honest MVP that runs; I’m not chasing perfection.”

This frames the day as **completion, not perfectionism**.

Phase 1 – Workspace & Cursor Setup (20–30 minutes)

Goal: Have everything ready in Cursor so you’re not configuring tools mid-sprint.

Steps:

1. **Open laptop, start Cursor, open the Guidelight repo**.
2. **Verify documentation is present:**

 - `README.md`
 - `GUIDELIGHT_SPEC.md`
 - `ARCHITECTURE_OVERVIEW.md`
 - `CONTRIBUTING.md`
 - `GUIDELIGHT_DEV_AGENT.md`

3. **Create / configure the Guidelight Dev Agent in Cursor:**

 - In Cursor’s agent/settings UI:
 - Add a new custom agent named something like ***“Guidelight Dev Agent”***.
 - Copy/paste the contents of `GUIDELIGHT_DEV_AGENT.md` into the agent’s system prompt.
 - Set **model** = `gpt-5.1` (or “GPT-5.1 Thinking”).

4. **Create a feature branch for the sprint:**

```
```bash
git checkout -b feature/guidelight-mvp
````
```

5. **Supabase environment sanity check:**

 - Ensure you have at least:
 - `VITE_SUPABASE_URL`
 - `VITE_SUPABASE_ANON_KEY`

- in your ` `.env.local` or appropriate config file.
- If anything is missing, note a small TODO like:
 - “TODO: fill Supabase env variables before data layer wiring.”

At this point, **you’re on a dedicated branch** and **your AI agent is configured** with the correct system prompt and model.

Phase 2 – Let GPT-5.1 Read & Plan (30–40 minutes)

Goal: Have the AI generate a concrete MVP plan rooted in your actual docs.

In Cursor (with Guidelight Dev Agent + GPT-5.1):

Use Composer/Agent and paste:

```text

You are the Guidelight Dev Agent.

Task:

1. Read and summarize briefly:
  - README.md
  - GUIDELIGHT\_SPEC.md
  - ARCHITECTURE\_OVERVIEW.md
2. Based on those documents, propose a concrete MVP plan that I can complete today:
  - 5–8 steps total.
  - Each step should be small enough to implement in a single Composer run.
  - For each step, list:
    - Goal.
    - Files/folders you expect to touch.
    - A suggested model (GPT-5.1 vs Claude Sonnet 4.5) using the “Model suggestion” line.
3. Do NOT modify any files yet. This is planning only.

```

What to do with the response:

- Read the summary to confirm the AI understood the project.
- Read the steps and make sure they roughly match what you expect:
 - Shell & routing,
 - Data layer,
 - Staff View,
 - Customer View,
 - Refactor / cleanup,
 - Review / docs.

If it looks mostly right, reply with:

```text

This plan looks good. Please restate the steps in a short, numbered list that I can reference

during the day, then stop. After that I will call you step by step.

---

Now you have a **\*\*roadmap created by the AI from your real docs\*\***.

If needed, copy the numbered list into a small file like `MVP\_PLAN\_TODAY.md` for quick reference.

---

## Phase 3 – App Shell & Routing (60–90 minutes)

**\*\*Goal:\*\*** Get the basic app running with routes and placeholder views.

**\*\*Model:\*\*** GPT-5.1 (default).

**\*\*Prompt example:\*\***

```text

Implement Step 1 of the MVP plan: app shell and routing.

Requirements:

- Vite + React + TypeScript setup (if not already done).
- A main layout and routing structure, with at least:
 - Customer View route.
 - Staff View route.
- Placeholder components for both views (they can show simple text for now).
- Follow the folder and file structure described in ARCHITECTURE_OVERVIEW.md.

Constraints:

- Do not add complex logic yet.
- Keep changes minimal and focused on getting the app to boot and route correctly.

At the end:

- List the files you created or modified.
- Confirm the npm scripts I should run to start the dev server.

****Then in terminal:****

```bash

```
npm install # if not already done
npm run dev

```

Open the app in the browser and confirm:

- It loads without errors.
- You can click or route to “Customer” and “Staff” views (even if they’re just placeholders).

\*\*Commit:\*\*

```
```bash
git add .
git commit -m "feat: scaffold Guidelight app shell and routing"
```
```

\*\*Mini break (5–10 minutes):\*\*

Stand up, stretch, maybe grab water. No doomscrolling.

---

## Phase 4 – Data Layer & Supabase Wiring (60–90 minutes)

\*\*Goal:\*\* Create a small, typed data layer that talks to Supabase and matches the schema.

\*\*Model:\*\* GPT-5.1.

\*\*Prompt example:\*\*

```text

Implement Step 2 of the MVP plan: data layer and Supabase wiring.

Goals:

- Use the Postgres MCP (if available) to inspect the Supabase dev schema.
- Based on GUIDELIGHT_SPEC.md and the actual schema, define TypeScript types for:
 - Budtenders.
 - Picks.
 - Any other absolutely essential entities.
- Implement a small API layer in src/lib/api (or as defined in ARCHITECTURE_OVERVIEW.md) with functions like:
 - listBudtenders()
 - listPicksForBudtender(budtenderId)
 - createPick(...)
 - updatePick(...)
 - deletePick(...)

Constraints:

- Do not put Supabase client calls directly in React components.
- Keep types aligned with the actual schema. If MCP is unavailable, describe assumptions clearly.

At the end:

- Summarize the schema you used.
- List the functions you created and their signatures.
- Note any assumptions or TODOs regarding the schema.

```

\*\*Then:\*\*

- Run `npm run build` or `npm run lint` to catch TypeScript issues.
- Fix any obvious errors with small follow-up prompts like:
  - “Fix the TS errors related to the new data layer in the simplest way that preserves type safety.”

\*\*Commit:\*\*

```
```bash
git add .
git commit -m "feat: add typed Supabase data layer for budtenders and picks"
```
```

\*\*Lunch break (30–45 minutes):\*\*

Eat something real. Walk a bit if you can. Let your head detach from the code for a bit.

---

## Phase 5 – Staff View MVP (90–120 minutes)

\*\*Goal:\*\* Staff can view and manage picks through the UI.

\*\*Model:\*\* GPT-5.1.

\*\*Prompt example:\*\*

```
```text
Implement Step 3 of the MVP plan: Staff View MVP.
```
```

Requirements from GUIDELIGHT\_SPEC.md:

- The Staff View should allow a staff member to:
  - See the list of budtender picks.
  - Create a new pick.
  - Edit an existing pick.
  - Delete or archive a pick.
- Use the data layer functions from src/lib/api; do not call Supabase directly from components.
- Handle loading, error, and empty states in a simple, honest way.
- Keep the layout clean and readable for backroom use; no need for fancy styling yet.

At the end:

- Describe the flow: what the staff member sees and how they interact.
- List which files were changed or created.
- Include a model suggestion line as configured in GUIDELIGHT\_DEV\_AGENT.md.

```

Then:

- Run the app and manually test:
 - Can you see picks?
 - Can you create/edit/delete a pick?
 - Does the UI show something sensible when there are no picks?

- Do errors show up in a non-confusing way?

Fix bugs as you find them—this may take a couple of small, targeted prompts.

****Commit:****

```
```bash
git add .
git commit -m "feat: implement Staff View MVP for managing picks"
```
```

****Short break (10–15 minutes):****

Stretch, maybe step away from the screen. Check in with how you're feeling. You're past the halfway point once Staff View works.

Phase 6 – Customer View MVP (60–90 minutes)

****Goal:**** Customer View displays picks in a clean, legible way.

****Model:**** GPT-5.1.

****Prompt example:****

```
```text
```

Implement Step 4 of the MVP plan: Customer View MVP.

Requirements:

- Display budtender picks in a way that is useful at the POS.
- For each pick, show:
  - Budtender identity (name or label).
  - Product/strain name.
  - 1–2 high-signal attributes (category, simple effect tags, or similar) as described in `GUIDELIGHT_SPEC.md`.
- Use the existing data layer; do not introduce new direct Supabase calls.
- Layout should be:
  - Readable from a short distance.
  - Not overloaded with text.
  - Grouped or organized in a way that makes sense for a budtender talking to a customer.

At the end:

- Describe how a budtender would actually use this view with a customer.
- Note any UX issues or future improvements you see but did not implement.

```

****Then:****

- Run the app and walk through it like:
 - A budtender at the counter, quickly referencing picks.

- A customer glancing at the screen.

Ask: "Is this **actually** helping the conversation?"

If something is obviously missing or confusing, fix the smallest possible thing that makes it better.

****Commit:****

```
```bash
git add .
git commit -m "feat: implement Customer View MVP for displaying budtender picks"
```
```

Phase 7 – Refactor & Cleanup with Claude (60–90 minutes)

****Goal:**** Use Claude's strengths once, in a controlled way, to improve clarity without chaos.

****Model:**** Switch to **Claude Sonnet 4.5** in Cursor for this phase only.

****Prompt example:****

```text

You are Claude Sonnet 4.5 acting as a refactor specialist for the Guidelight app.

Context:

- Guidelight is a Vite + React + TypeScript app using Supabase for data.
- We now have working Staff and Customer views wired through src/lib/api.
- Please read:
  - GUIDELIGHT\_SPEC.md
  - ARCHITECTURE\_OVERVIEW.md
  - GUIDELIGHT\_DEV\_AGENT.md

Task:

1. Inspect:

- src/views/StaffView\*.tsx
- src/views/CustomerView\*.tsx
- src/lib/api/\*

2. Propose a SMALL set of improvements that will:

- Simplify data and prop flow.
- Improve type safety.
- Remove obvious duplication.

3. Apply these improvements in a scoped way:

- Do NOT rewrite the entire app or change the high-level structure.
- Limit changes to these files and any minimal helper modules they depend on.

4. Summarize:

- What you changed.
- Why the code is now clearer/safer.

Keep the refactor surgical, not sweeping.

---

\*\*Your job here:\*\*

- Review diffs carefully. If Claude tries to change too much, tell it explicitly:

- "That's too large a refactor. Revert X and keep Y, focusing only on simplifying data flow in StaffView for now."

- Re-run the app and basic checks:

- `npm run dev`

- `npm run build` or `npm run lint`

\*\*Commit:\*\*

```
```bash
git add .
git commit -m "refactor: clarify views and data layer for Guidelight MVP"
````
```

---

## Phase 8 – Codex Review & Docs Sync (60 minutes)

\*\*Goal:\*\* Let an external brain (Codex) review your work, then align docs with reality.

### 8.1 Push branch & open PR

```
```bash
git push -u origin feature/guidelight-mvp
````
```

Open a PR on GitHub from `feature/guidelight-mvp` → `main`.

### 8.2 Codex IDE review

In Codex IDE, give it the PR or diff and say something like:

```text

This is the Guidelight MVP pull request.

Please review the diff with focus on:

- TypeScript correctness and type safety.
- Potential runtime errors (null/undefined, missing fields, unhandled promises).
- Data-fetching edge cases (loading, error, empty states).
- Any obviously over-complicated logic that could be simplified.

Output:

- A numbered list of concrete issues with file + line hints.
- Plain-language suggestions for fixes (no full rewrites).
- Any high-level design concerns that are worth noting but not required to block this MVP.

``

8.3 Apply high-value fixes with GPT-5.1 in Cursor

Back in Cursor (switch model back to GPT-5.1), ask it to apply Codex's best suggestions, one cluster at a time. For example:

```text

Here is Codex's review of the Guidelight MVP PR (paste excerpt).

Please:

1. Fix the concrete issues they've listed related to null/undefined handling in StaffView and CustomerView.
2. Keep changes minimal and focused.
3. Summarize what you changed.

Do not restructure anything beyond what is required for these fixes.

``

Re-run the app and basic checks after changes.

\*\*Commit:\*\*

```bash

```
git add .
git commit -m "chore: address Codex review feedback for Guidelight MVP"
```

``

8.4 Update docs with GPT-5.1

Ask GPT-5.1 in Cursor to bring docs up to date:

```text

Update the documentation to reflect the current Guidelight MVP:

1. Update README.md to:
  - Briefly describe the current MVP features (Staff View and Customer View).
  - Mention that this is an internal tool for State of Mind.
2. Review GUIDELIGHT\_SPEC.md and adjust any details that are clearly outdated compared to the implementation.
3. If the architecture differs meaningfully from ARCHITECTURE\_OVERVIEW.md, suggest small edits to align it (or add a short "Current Status" note).

Make the smallest, clearest edits that make the docs honest about the MVP.

``

Review the changes, tweak anything that feels off, and commit:

```bash

```
git add .
```

```
git commit -m "docs: sync Guidelight docs with MVP implementation"
```

```
---
```

Phase 9 – Wrap-Up & Decompression (15–30 minutes)

Goal: End the day with clarity instead of mental static.

1. **Create or update `NEXT_STEPS.md` (or similar)**

- Bullet out:

- Small feature gaps (e.g., better filters, sorting, metadata).
- Known tech debt.
- Any future ideas you don't want to lose.

2. **Short personal note (optional but helpful)**

- One paragraph somewhere (Obsidian, Notion, a text file):

- What you actually got done.
- What you're proud of.
- What felt hard or confusing.

3. **Step away**

- Close Cursor.
- Do something off-screen for a while (music, show, walk, etc.).
- Let tomorrow's work wait.

```
---
```

4. Quick-Reference Prompts

You can copy these snippets into a small cheat sheet file if you want.

4.1 Plan the day (GPT-5.1)

```
```text
```

Read README.md, GUIDELIGHT\_SPEC.md, and ARCHITECTURE\_OVERVIEW.md.

Then:

- Summarize the project in a few sentences.
- Propose a 5–8 step MVP plan for today with goals and files per step.
- Add a “Model suggestion” line for each step based on our rules.

Do not edit files yet.

```
```
```

4.2 Implement a step (GPT-5.1)

```
```text
```

Implement Step X of today's MVP plan.

- Focus only on this step.
- Minimize unrelated changes.
- Use existing patterns and the data layer.
- At the end, list files changed and summarize behavior.

Model suggestion: include it, but proceed with GPT-5.1.

---

### 4.3 Claude refactor

```text

You are Claude Sonnet 4.5 acting as a refactor specialist.

Inspect:

- src/views/StaffView*.tsx
- src/views/CustomerView*.tsx
- src/lib/api/*

Propose and apply a small set of improvements to simplify data/prop flow and improve type safety without changing overall behavior or structure. Keep changes surgical.

4.4 Codex review

```text

Review this Guidelight MVP pull request.

Focus on:

- TS correctness and type safety.
- Null/undefined and runtime risks.
- Data-fetch edge cases.
- Overly complex logic.

Return a numbered list of issues with file/line hints and plain-language fix suggestions.

---

---

## 5. Closing Note

This plan is not a rigid contract; it's a \*\*scaffold\*\*. If tomorrow goes differently (longer data-layer work, faster UI work, etc.), that's fine. The important part is:

- You start with a clear \*\*stack and model strategy\*\*.
- You let the AI do the heavy lifting where it's strongest.
- You stay in control of architecture, UX, and quality.
- You end the day with a \*\*real, running Guidelight MVP\*\* and a sense of where to go next.

Use this document however feels natural: read it through once tonight, then tomorrow just dip into the relevant sections as you move through the day.